

✓ American Sign Language Recognition Project

By Jacob Flores & Jason Flores

First we imported all packages needed, then we installed kaggle in order to download the American Sign Language dataset. Before exploring the data we confirmed that the amount of classes are the correct (29). Also, we made sure to have latest version of PyTorch and Torchvision.

✓ 0 - Imports

```
1 import torch
2 from torch import nn
3 import torch.nn as nn
4 import torch.optim as optim
5 import torchvision
6 from torch.utils.data import DataLoader, Subset, random_split
7 from torchvision import datasets, transforms, models
8 from torchvision.transforms import ToTensor
9 import matplotlib.pyplot as plt
10 import numpy as np
11 import random
12 import os
13
14 !pip install -q kaggle
15
16 os.environ['KAGGLE_USERNAME'] = "jacobbb593"
17 os.environ['KAGGLE_KEY'] = "56bc160421849119bc1cfde80d32e83f"
18
19 !kaggle datasets download -d grassknotted/asl-alphabet
20
21 !unzip -q -o asl-alphabet.zip -d asl_alphabet_data
22
23
24 data_path = "/content/asl_alphabet_data/asl_alphabet_train/asl_alphabet_train"
25 print("Classes found:", os.listdir(data_path)[:5])
26 print("Total classes:", len(os.listdir(data_path)))
27
28 # Check Version
29 print(f"PyTorch version: {torch.__version__}\ntorchvision version: {torchvision.__version__}")
```



Dataset URL: <https://www.kaggle.com/datasets/grassknotted/asl-alphabet>
 License(s): GPL-2.0
 asl-alphabet.zip: Skipping, found more recently modified local copy (use --force to force download)
 Classes found: ['R', 'V', 'H', 'T', 'Y']
 Total classes: 29
 PyTorch version: 2.6.0+cu124
 torchvision version: 0.21.0+cu124

✓ 1 - Getting a Dataset

After having the dataset download we decided to modified it based on our needs. So we transformed it by reducing the size of images from 200x200(originally) to 64x64 so that we can process the images faster.

After having that modified, we decided to reduce our dataset (from 87,000 images to 5,000 images). This is because we are looking for efficiency in our project and to reduce the amount of time during the training.

As usual, we split into 80% train and 20% test so that we have a balance between the data.

After having that defined we create a Dataloader for training and test set. Then we print the dataset that we have defined.

After having the dataset that we are going to use, we explored the whole dataset by looking at lenghts, tensor values, classes names, classes index. In addition we checked input and output shapes of our data and then we visualize each class.

```
1 data_dir = "/content/asl_alphabet_data/asl_alphabet_train/asl_alphabet_train"

1 transform = transforms.Compose([
2     transforms.Resize((64, 64)),
```

```

3     transforms.ToTensor(),
4     transforms.Normalize((0.5,),(0.5,))
5 ])
6
7 full_dataset = datasets.ImageFolder(root=data_dir,
8                                     transform=transform)
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60
61
62
63
64
65
66
67
68
69
70
71
72
73
74
75
76
77
78
79
80
81
82
83
84
85
86
87
88
89
90
91
92
93
94
95
96
97
98
99
100
101
102
103
104
105
106
107
108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124
125
126
127
128
129
130
131
132
133
134
135
136
137
138
139
140
141
142
143
144
145
146
147
148
149
150
151
152
153
154
155
156
157
158
159
160
161
162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781
782
783
784
785
786
787
788
789
790
791
792
793
794
795
796
797
798
799
800
801
802
803
804
805
806
807
808
809
810
811
812
813
814
815
816
817
818
819
820
821
822
823
824
825
826
827
828
829
830
831
832
833
834
835
836
837
838
839
840
841
842
843
844
845
846
847
848
849
850
851
852
853
854
855
856
857
858
859
860
861
862
863
864
865
866
867
868
869
870
871
872
873
874
875
876
877
878
879
880
881
882
883
884
885
886
887
888
889
890
891
892
893
894
895
896
897
898
899
900
901
902
903
904
905
906
907
908
909
910
911
912
913
914
915
916
917
918
919
920
921
922
923
924
925
926
927
928
929
930
931
932
933
934
935
936
937
938
939
940
941
942
943
944
945
946
947
948
949
950
951
952
953
954
955
956
957
958
959
960
961
962
963
964
965
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988
989
990
991
992
993
994
995
996
997
998
999

```

Using 16 samples from 29 classes

```
1 len(train_data.dataset), len(test_data.dataset)
```

(5000, 5000)

```
1 len(train_data), len(test_data)
```

(4000, 1000)

```
1 image, label = train_data[0]
```

```
2 image, label
```

```

(tensor([[-0.8275, -0.6314, -0.6314, ..., -0.7569, -0.6784, -0.6706],
        [-0.6471, -0.2157, -0.1922, ..., -0.2941, -0.0980, -0.2627],
        [-0.6941, -0.3098, -0.2941, ..., -0.0196, -0.0275, -0.2784],
        ...,
        [-0.4275, 0.3961, 0.4039, ..., 0.1451, 0.1765, -0.1137],
        [-0.4275, 0.3725, 0.4039, ..., 0.1373, 0.1922, -0.1137],
        [-0.5686, 0.0353, 0.0510, ..., -0.1451, -0.1294, -0.3569]],
        [[-0.8431, -0.6549, -0.6392, ..., -0.7961, -0.7098, -0.7098],
        [-0.7020, -0.2863, -0.2235, ..., -0.4039, -0.1608, -0.2941],
        [-0.7176, -0.3569, -0.3725, ..., -0.1294, -0.0824, -0.3020],
        ...,
        [-0.3725, 0.5059, 0.5137, ..., 0.0902, 0.1294, -0.1216],
        [-0.3647, 0.5137, 0.5059, ..., 0.0745, 0.1373, -0.1137],
        [-0.5216, 0.1686, 0.1529, ..., -0.1608, -0.1294, -0.3255]],
        [[ 0.6784, 0.3569, 0.3412, ..., 0.2627, 0.3412, 0.5922],
        [ 0.3333, -0.4039, -0.4196, ..., -0.4510, -0.1686, 0.1843],
        [ 0.3020, -0.4745, -0.5451, ..., -0.1922, -0.0824, 0.1843],
        ...,
        [ 0.7490, 0.5686, 0.5059, ..., 0.1059, 0.1373, 0.3961],
        [ 0.7412, 0.5608, 0.5137, ..., 0.0745, 0.1137, 0.3725],
        [ 0.7647, 0.6078, 0.6000, ..., 0.3098, 0.3255, 0.4902]]]),
9)

```

```
1 class_names = full_dataset.classes
```

```
2 class_names
```

['A',
'B',
'C',

```
'D',
'E',
'F',
'G',
'H',
'I',
'J',
'K',
'L',
'M',
'N',
'O',
'P',
'Q',
'R',
'S',
'T',
'U',
'V',
'W',
'X',
'Y',
'Z',
'del',
'nothing',
'space']
```

```
1 class_to_idx = full_dataset.class_to_idx
2 class_to_idx
```

```
{'A': 0,
'B': 1,
'C': 2,
'D': 3,
'E': 4,
'F': 5,
'G': 6,
'H': 7,
'I': 8,
'J': 9,
'K': 10,
'L': 11,
'M': 12,
'N': 13,
'O': 14,
'P': 15,
'Q': 16,
'R': 17,
'S': 18,
'T': 19,
'U': 20,
'V': 21,
'W': 22,
'X': 23,
'Y': 24,
'Z': 25,
'del': 26,
'nothing': 27,
'space': 28}
```

✓ 1.1 - Check input and output shapes of data

```
1 # Check the shape of our image
2 print(f"Image shape: {image.shape} -> [color_channels, height, width]")
3 print(f"Image label: {class_names[label]}")
```

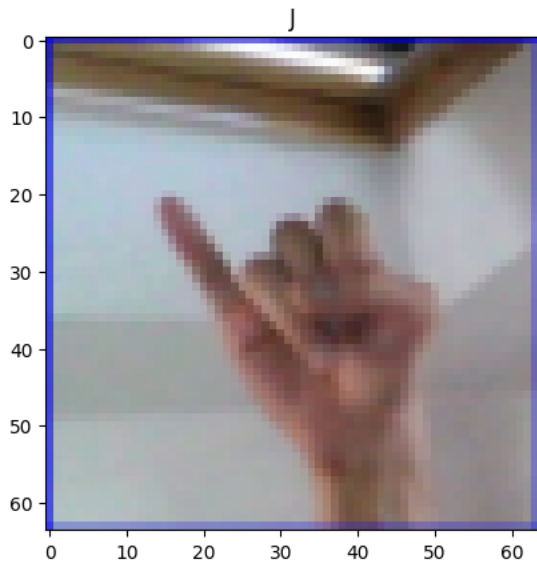
```
Image shape: torch.Size([3, 64, 64]) -> [color_channels, height, width]
Image label: J
```

✓ 1.2 - Visualizing our data

```
1 # Before we show a sample, notice how we change the image values, this is because
2 # we are performing reverse normalization in order to have pixel values in the range and avoid errors
3 image = image * 0.5 + 0.5 # You are going to see this more often along the way
4 print(f"Image shape: {image.shape}")
```

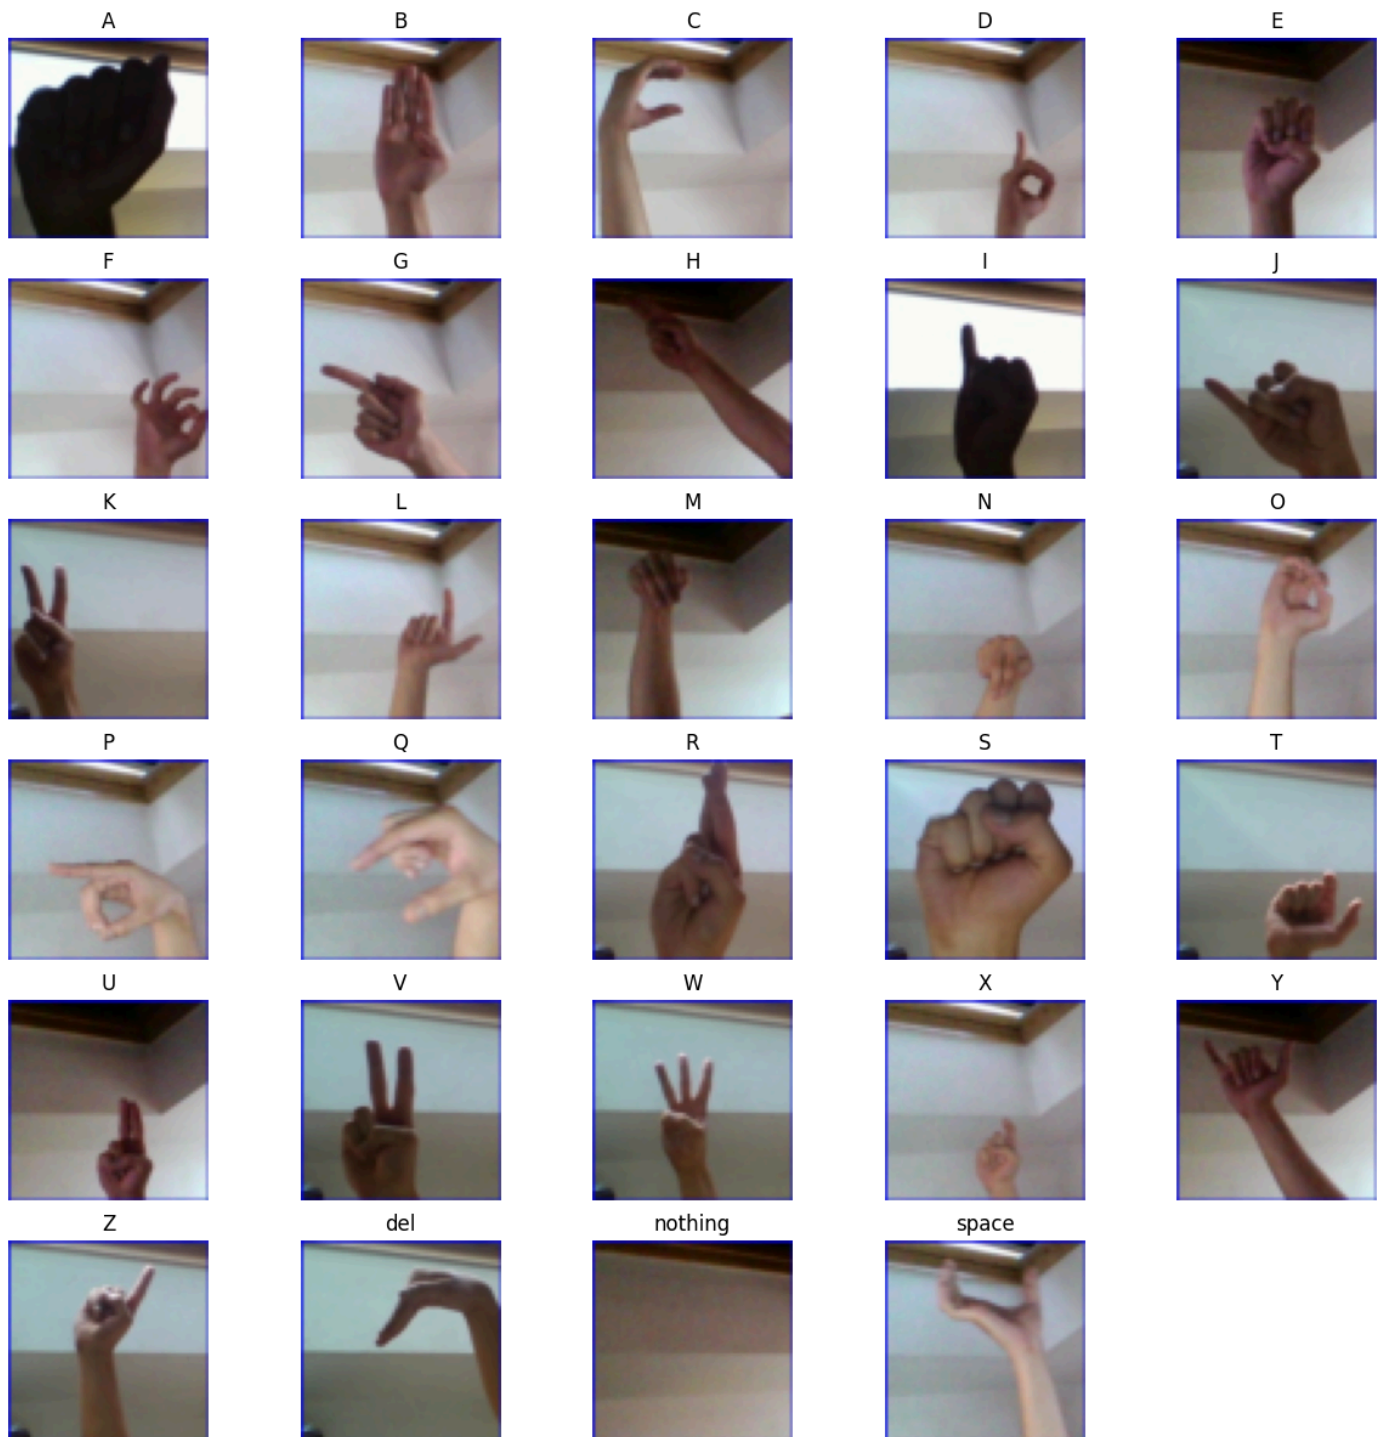
```
5 plt.imshow(image.permute(1, 2, 0))
6 plt.title(classes[label]);
```

Image shape: torch.Size([3, 64, 64])



```
1 # Get one image for each class
2 images = []
3 labels = []
4 for label in range(len(classes)):
5     # Find the index of the first image with the current label
6     index = next((i for i, (image, target) in enumerate(subset) if target == label), None)
7     if index is not None:
8         image, target = subset[index]
9         images.append(image)
10        labels.append(target)
```

```
1 # Plot images
2 fig = plt.figure(figsize=(15, 15))
3 rows, cols = 6, 5 # Adjust rows and cols to fit all classes
4 for i in range(len(images)):
5     img, label = images[i], labels[i]
6
7     # Reverse normalization for display as mentioned earlier
8     img = img * 0.5 + 0.5
9
10    fig.add_subplot(rows, cols, i + 1)
11    plt.imshow(img.permute(1, 2, 0))
12    plt.title(class_names[label])
13    plt.axis(False)
14
```



```
1 train_data, test_data
```

```
<torch.utils.data.dataset.Subset at 0x7e1fb489ec90>,  
<torch.utils.data.dataset.Subset at 0x7e1fb489e950>)
```

2 - Prepare DataLoader

This part prepares data batches for training and testing. First we defined our batch size. By using DataLoader, the training and testing datasets are converted into iterable objects. As we can see in the code the training dataloader shuffles the data in each epoch so that we

prevent the model to learn the order of the data. But the testing Dataloader does not shuffle the data as we rather consistent evaluation during the testing.

Then we check out the output, and further inspect the data by retrieving a single batch and printing its shape to confirm it is in the format we expect.

Finally an image from the batch is shown along with its label and shape to ensure the data is loaded and labeled correctly. This process is important to make sure that we are in the right track.

```
1 # Setup the batch size hyperparameter
2 BATCH_SIZE = 32
3
4 # Turn datasets into iterables (batches)
5 train_dataloader = DataLoader(train_data, # dataset to turn into iterable
6                               batch_size=BATCH_SIZE, # how many samples per batch?
7                               shuffle=True) # shuffle data every epoch?
8
9 test_dataloader = DataLoader(test_data,
10                             batch_size=BATCH_SIZE,
11                             shuffle=False) # don't necessarily have to shuffle the testing data
12
13 train_dataloader, test_dataloader
```

```
(<torch.utils.data.dataloader.DataLoader at 0x7e1fb437c9d0>,
 <torch.utils.data.dataloader.DataLoader at 0x7e1fb4158c10>)
```

```
1 # Let's check out what we've created
2 print(f"Dataloaders: {train_dataloader, test_dataloader}")
3 print(f"Length of train dataloader: {len(train_dataloader)} batches of {BATCH_SIZE}")
4 print(f"Length of test dataloader: {len(test_dataloader)} batches of {BATCH_SIZE}")
```

```
Dataloaders: (<torch.utils.data.dataloader.DataLoader object at 0x7e1fb437c9d0>, <torch.utils.data.dataloader.DataLoader obj
Length of train dataloader: 125 batches of 32
Length of test dataloader: 32 batches of 32
```

```
1 # Check out what's inside the training dataloader
2 train_features_batch, train_labels_batch = next(iter(train_dataloader))
3 train_features_batch.shape, train_labels_batch.shape
```

```
(torch.Size([32, 3, 64, 64]), torch.Size([32]))
```

```
1 # Show a sample
2 torch.manual_seed(42)
3 random_idx = torch.randint(0, len(train_features_batch), size=[1]).item()
4 img, label = train_features_batch[random_idx], train_labels_batch[random_idx]
5
6 # Reverse normalization for display
7 img = img * 0.5 + 0.5
8
9 plt.imshow(img.permute(1, 2, 0))
10 plt.title(class_names[label])
11 plt.axis(False);
12 print(f"Image size: {img.shape}")
13 print(f"Label: {label}, label size: {label.shape}")
```

Image size: torch.Size([3, 64, 64])
 Label: 9, label size: torch.Size([1])

J



✓ 3 - BASELINE MODEL

After having worked on our dataset and dataloader, for the sake of getting to know our model better, we decided to implement a baseline model where we are using a flatten layer and where we are evaluating a simple sample. We implemented it because is the best way to know our model needs depending on how high the accuracy is or how low the loss is. Later on, we'll be creating more a non-linear model and a CNN model to see how much our model is learning in each one of them.

For making complex things we first have to start simple.

```

1 # Create a flatten layer
2 flatten_model = nn.Flatten()
3
4 # Get a single sample
5 x = train_features_batch[0]
6
7 # Flatten the sample
8 output = flatten_model(x) # perform forward pass
9
10 # Print out what happened
11 print(f"Shape before flattening: {x.shape} -> [color_channels, height, width]")
12 print(f"Shape after flattening: {output.shape} -> [color_channels, height*width]")

```

Shape before flattening: torch.Size([3, 64, 64]) -> [color_channels, height, width]
 Shape after flattening: torch.Size([3, 4096]) -> [color_channels, height*width]

```

1 from torch import nn
2
3 # Baseline model definition along with its forward function
4 # we did it with just two linear layers to make it simple enough
5 class ASLBaseline(nn.Module):
6     def __init__(self, input_shape: int, hidden_units: int, output_shape: int):
7         super().__init__()
8         self.layer_stack = nn.Sequential(
9             nn.Flatten(),
10             nn.Linear(in_features=input_shape,
11                       out_features=hidden_units),
12             nn.Linear(in_features=hidden_units,
13                       out_features=output_shape)
14         )
15
16     def forward(self, x):
17         return self.layer_stack(x)

```

```

1 torch.manual_seed(42)
2
3 # Create an instance model_0

```

```

4 model_0 = ASLBaseline(
5     input_shape = 3*64*64,
6     hidden_units = 10,
7     output_shape=len(class_names)
8 ).to("cpu")
9
10 model_0

ASLBaseline(
  (layer_stack): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=12288, out_features=10, bias=True)
    (2): Linear(in_features=10, out_features=29, bias=True)
  )
)

1 model_0.state_dict()

OrderedDict([('layer_stack.1.weight',
  tensor([[ 6.8970e-03,  7.4876e-03, -2.1134e-03, ...,  6.6472e-04,
           -4.7116e-03, -5.2405e-03],
          [-1.9907e-04,  3.1398e-03, -8.3313e-03, ...,  4.0213e-03,
           6.6041e-03,  5.0459e-03],
          [ 6.0608e-03, -5.4492e-03,  7.6998e-03, ...,  3.5302e-05,
           -6.7376e-03,  7.3019e-03],
          ...,
          [-4.3874e-03,  1.9926e-04,  7.9915e-03, ..., -5.0085e-03,
           -3.9315e-03, -7.5393e-03],
          [ 4.2025e-03, -1.3148e-03, -2.7806e-03, ..., -5.4377e-03,
           -4.8458e-03,  3.0181e-03],
          [ 3.0575e-03,  5.3473e-03, -7.9494e-03, ..., -8.9380e-03,
           8.5260e-03,  2.1958e-03]])),
  ('layer_stack.1.bias',
  tensor([ 0.0036,  0.0022,  0.0069, -0.0080, -0.0077,  0.0016, -0.0080,  0.0013,
           0.0007,  0.0080])),
  ('layer_stack.2.weight',
  tensor([[ 0.2271, -0.2909,  0.0526,  0.2420, -0.0226,  0.1177,  0.0311,  0.0199,
           0.0727,  0.0551],
          [-0.1879, -0.1115, -0.2149,  0.1204, -0.2628,  0.0700,  0.0217, -0.1841,
           -0.1254, -0.0144],
          [-0.1534,  0.1803, -0.1216, -0.2853,  0.1073,  0.1263, -0.2445,  0.1226,
           0.1612, -0.2763],
          [ 0.1309,  0.0899, -0.1710, -0.1101, -0.1523,  0.1366, -0.1008, -0.0388,
           -0.1357, -0.1274],
          [ 0.2154, -0.0421,  0.0956, -0.0118, -0.2286,  0.3034,  0.2304, -0.1224,
           0.1647,  0.0040],
          [ 0.1699, -0.0286,  0.1463, -0.2019, -0.1383, -0.0608, -0.0046,  0.1950,
           0.2775, -0.1048],
          [-0.2643, -0.2789,  0.2549,  0.1814, -0.2726,  0.0818, -0.0546, -0.2327,
           0.1538,  0.3080],
          [-0.2690,  0.2196,  0.2174, -0.2745,  0.1323, -0.0436, -0.0068, -0.1576,
           0.0625,  0.0813],
          [ 0.1939,  0.1772, -0.2130,  0.2027, -0.0021, -0.1484, -0.2197, -0.0111,
           0.0386,  0.0530],
          [-0.2387, -0.0817,  0.2055,  0.0559,  0.2870, -0.3095, -0.1049, -0.0042,
           -0.1549,  0.3149],
          [ 0.2623,  0.0288,  0.1195, -0.2930, -0.0005,  0.3093,  0.1367,  0.2194,
           0.0865,  0.0384],
          [ 0.0579,  0.1749, -0.1450, -0.0794,  0.2511, -0.1864, -0.1222,  0.1361,
           -0.1253,  0.2578],
          [-0.3055, -0.1784,  0.1948, -0.2012, -0.1424, -0.1791,  0.2238, -0.0434,
           -0.1099,  0.2489],
          [-0.1836, -0.1452,  0.3038,  0.0538,  0.0413,  0.2980, -0.0539,  0.2846,
           0.1856, -0.0064],
          [ 0.0710, -0.1581,  0.1057, -0.0919,  0.1336, -0.2215,  0.1867, -0.2535,
           -0.0709,  0.1048],
          [ 0.2143,  0.2477, -0.0449,  0.2926, -0.0164,  0.2143, -0.1815,  0.1759,
           0.0603,  0.0334],
          [ 0.0896, -0.1494,  0.2653, -0.0646, -0.2180,  0.0760, -0.0473,  0.1632,
           -0.0911, -0.0837],
          [-0.2410,  0.1765, -0.1889, -0.0837,  0.2526, -0.1805, -0.2914, -0.0554,
           -0.0757,  0.1784],
          [ 0.2267,  0.1200,  0.2232, -0.2571, -0.1645, -0.1279,  0.0547, -0.1862,
           -0.0088,  0.0695],
          [-0.2006,  0.1453, -0.1304, -0.0065,  0.2275,  0.1927, -0.1856,  0.0585,
           -0.0081, -0.0397]]))
])

```

- 3.1 - Setup loss, optimizer and evaluation metrics


```

1 import requests
2 from pathlib import Path
3
4 # Download helper functions if we do not have it. This will help us having an accuracy value
5 if Path("helper_functions.py").is_file():
6     print("helper_functions.py already exists, skipping download...")
7 else:
8     print("Downloading helper_functions.py")
9     request = requests.get("https://raw.githubusercontent.com/mrdbourke/pytorch-deep-learning/main/helper_functions.py")
10    with open("helper_functions.py", "wb") as f:
11        f.write(request.content)

```

↗ helper_functions.py already exists, skipping download...

```

1 # Import accuracy metric
2 from helper_functions import accuracy_fn
3
4 # Setup loss function and optimizer
5 loss_fn = nn.CrossEntropyLoss() # For multi classification
6 optimizer = torch.optim.SGD(params=model_0.parameters(), # stochastic gradient descent as usual
7                               lr=0.01)

```

✓ 3.2 - Creating a function to time our experiments

As ML is very experimental, two of the main things we'll often want to track are the model's performance and how fast it runs

Please find our timing function below:

```

1 from timeit import default_timer as timer
2
3 def print_train_time(start: float,
4                      end: float,
5                      device: torch.device = None):
6     """Prints difference between start and end time."""
7     total_time = end - start
8     print(f"Train time on {device}: {total_time:.3f} seconds")
9     return total_time

```

```

1 # Quick example to see how if it works
2 start_time = timer()
3
4 end_time = timer()
5 print_train_time(start=start_time, end=end_time, device="cpu")

```

↗ Train time on cpu: 0.000 seconds
4.5781999915561755e-05

✓ 3.3 - Creating a training loop and training model on batches of data

```

1 # Import tqdm for progress bar
2 from tqdm.auto import tqdm
3
4 # Set the seed and start the timer
5 torch.manual_seed(42)
6 train_time_start_on_cpu = timer()
7
8 # Set the number of epochs (we'll keep this small for faster training time as mentioned in our first presentation)
9 epochs = 3
10
11 # Create training and test loop
12 for epoch in tqdm(range(epochs)):
13     print(f"Epoch: {epoch}\n-----")
14     ### Training
15     train_loss = 0
16     # Add a loop to loop through the training batches
17     for batch, (X, y) in enumerate(train_data_loader):
18         model_0.train()
19         # 1. Forward pass
20         y_pred = model_0(X)
21

```

```

22 # 2. Calculate loss (per batch)
23 loss = loss_fn(y_pred, y)
24 train_loss += loss # accumulate train loss
25
26 # 3. Optimizer zero grad
27 optimizer.zero_grad()
28
29 # 4. Loss backward
30 loss.backward()
31
32 # 5. Optimizer step (update the model's parameters once *per batch*)
33 optimizer.step()
34
35 # Print out what's happening
36 if batch % 400 == 0:
37     print(f"Looked at {batch * len(X)}/{len(train_dataloader.dataset)} samples.")
38
39 # Divide total train loss by length of train dataloader
40 train_loss /= len(train_dataloader)
41
42 ### Testing
43 test_loss, test_acc = 0, 0
44 model_0.eval()
45 with torch.inference_mode():
46     for X_test, y_test in test_dataloader:
47         # 1. Forward pass
48         test_pred = model_0(X_test)
49
50         # 2. Calculate loss (accumulatively)
51         test_loss += loss_fn(test_pred, y_test)
52
53         # 3. Calculate accuracy
54         test_acc += accuracy_fn(y_true=y_test, y_pred=test_pred.argmax(dim=1))
55
56 # Calculate the test loss average per batch
57 test_loss /= len(test_dataloader)
58
59 # Calculate the test acc average per batch
60 test_acc /= len(test_dataloader)
61
62 # Print out what's happening
63 print(f"\nTrain loss: {train_loss:.4f} | Test loss: {test_loss:.4f}, Test acc: {test_acc:.4f}")
64
65 # Calculate training time
66 train_time_end_on_cpu = timer()
67 total_train_time_model_0 = print_train_time(start=train_time_start_on_cpu,
68                                             end=train_time_end_on_cpu,
69                                             device=str(next(model_0.parameters()).device))

```

100% 3/3 [00:21<00:00, 6.90s/it]

```

Epoch: 0
-----
Looked at 0/4000 samples.

Train loss: 3.2142 | Test loss: 3.0513, Test acc: 16.1133
Epoch: 1
-----
Looked at 0/4000 samples.

Train loss: 2.9499 | Test loss: 2.9235, Test acc: 18.3594
Epoch: 2
-----
Looked at 0/4000 samples.

Train loss: 2.7644 | Test loss: 2.8075, Test acc: 19.4336
Train time on cpu: 21.860 seconds

```

✓ 3.4 - Make predictions and get Model 0 results

```

1 torch.manual_seed(42)
2 def eval_model(model: torch.nn.Module,
3               data_loader: torch.utils.data.DataLoader,
4               loss_fn: torch.nn.Module,
5               accuracy_fn):
6     """Returns a dictionary containing the results of model predicting on data_loader."""

```

```

7  loss, acc = 0, 0
8  model.eval()
9  with torch.inference_mode():
10     for X, y in tqdm(data_loader):
11         # Make predictions
12         y_pred = model(X)
13
14         # Accumulate the loss and acc values per batch
15         loss += loss_fn(y_pred, y)
16         acc += accuracy_fn(y_true=y,
17                           y_pred=y_pred.argmax(dim=1))
18
19     # Scale loss and acc to find the average loss/acc per batch
20     loss /= len(data_loader)
21     acc /= len(data_loader)
22
23     return {"model_name": model.__class__.__name__, # only works when model was created with a class
24           "model_loss": loss.item(),
25           "model_acc": acc}
26
27 # Calculate model 0 results on test dataset
28 model_0_results = eval_model(model=model_0,
29                               data_loader=test_data_loader,
30                               loss_fn=loss_fn,
31                               accuracy_fn=accuracy_fn)
32 model_0_results

```

100% 32/32 [00:01<00:00, 21.88it/s]

```

{'model_name': 'ASLBaseline',
 'model_loss': 2.8074750900268555,
 'model_acc': 19.43359375}

```

3.5 Setup device agnostic-code (for using a GPU if there is one)

```
1 !nvidia-smi
```

Fri Apr 25 17:10:38 2025

NVIDIA-SMI 550.54.15				Driver Version: 550.54.15		CUDA Version: 12.4		
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.	
0	Tesla T4	Off	00000000:00:04.0	Off	0			
N/A	47C	P8	9W / 70W	2MiB / 15360MiB	0%	Default	N/A	

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
ID	ID					Usage	
No running processes found							

```
1 torch.cuda.is_available()
```

True

```

1 # Setup device-agnostic code
2 import torch
3 device = "cuda" if torch.cuda.is_available() else "cpu"
4 device

```

'cuda'

4 - Model 1: Building a better model with non-linearity

Once built our baseline model, we've noticed the lack of accuracy as well as a high loss, reason why we've decided to implement a better model with non-linearity.

```

1 # Create a model with non-linear and linear layers along with its forward function
2 # Here we are using ReLU unlike the baseline model where we treated it as a linear model
3 class ASLModelV1(nn.Module):
4     def __init__(self,
5                   input_shape: int,
6                   hidden_units: int,
7                   output_shape: int):
8         super().__init__()
9         self.layer_stack = nn.Sequential(
10             nn.Flatten(), # flatten inputs into a single vector
11             nn.Linear(in_features=input_shape,
12                      out_features=hidden_units),
13             nn.ReLU(),
14             nn.Linear(in_features=hidden_units,
15                      out_features=output_shape),
16             nn.ReLU()
17         )
18
19     def forward(self, x: torch.Tensor):
20         return self.layer_stack(x)

1 # Create an instance of model_1
2 torch.manual_seed(42)
3 model_1 = ASLModelV1(input_shape=12288, # this is the output of the flatten after our 3*64*64 image goes in
4                      hidden_units=10,
5                      output_shape=len(class_names)).to(device) # send to the GPU if it's available
6 next(model_1.parameters()).device

device(type='cuda', index=0)

```

✓ 4.1 - Setup loss, optimizer and evaluation metrics

```

1 from helper_functions import accuracy_fn
2 loss_fn = nn.CrossEntropyLoss() # measure how wrong our model is
3 optimizer = torch.optim.SGD(params=model_1.parameters(), # tries to update our model's parameters to reduce the loss
4                             lr=0.01)

```

✓ 4.2 - Functoinizing training and evaluation/testing loops

```

1 # Here we are defining train_step in order to avoid rewriting it over and over again in case we implement more models
2 # in that way we would be saving memory and time, and the process would run smoothly and be more efficient
3 def train_step(model: torch.nn.Module,
4               data_loader: torch.utils.data.DataLoader,
5               loss_fn: torch.nn.Module,
6               optimizer: torch.optim.Optimizer,
7               accuracy_fn,
8               device: torch.device = device):
9     """Performs a training with model trying to learn on data_loader."""
10    train_loss, train_acc = 0, 0
11
12    # Put model into training mode
13    model.train()
14
15    # Add a loop to loop through the training batches
16    for batch, (X, y) in enumerate(data_loader):
17        # Put data on target device
18        X, y = X.to(device), y.to(device)
19
20        # 1. Forward pass (outputs the raw logits from the model)
21        y_pred = model(X)
22
23        # 2. Calculate loss and accuracy (per batch)
24        loss = loss_fn(y_pred, y)
25        train_loss += loss # accumulate train loss
26        train_acc += accuracy_fn(y_true=y,
27                                y_pred=y_pred.argmax(dim=1)) # go from logits -> prediction labels
28

```

```

29 # 3. Optimizer zero grad
30 optimizer.zero_grad()
31
32 # 4. Loss backward
33 loss.backward()
34
35 # 5. Optimizer step (update the model's parameters once *per batch*)
36 optimizer.step()
37
38 # Divide total train loss and acc by length of train dataloader
39 train_loss /= len(data_loader)
40 train_acc /= len(data_loader)
41 print(f"Train loss: {train_loss:.5f} | Train acc: {train_acc:.2f}%")
42
43 # return the loss value and accuracy for future linear plots
44 return train_loss.item(), train_acc

1 # Here we are defining test_step in order to avoid rewriting it over and over again in case we implement more models
2 # in that way we would be saving memory and time, and the process would run smoothly and be more efficient
3 def test_step(model: torch.nn.Module,
4               data_loader: torch.utils.data.DataLoader,
5               loss_fn: torch.nn.Module,
6               accuracy_fn,
7               device: torch.device = device):
8     """Performs a testing loop step on model going over data_loader."""
9     test_loss, test_acc = 0, 0
10
11 # Put the model in eval mode
12 model.eval()
13
14 # Turn on inference mode context manager
15 with torch.inference_mode():
16     for X, y in data_loader:
17         # Send the data to the target device
18         X, y = X.to(device), y.to(device)
19
20         # 1. Forward pass (outputs raw logits)
21         test_pred = model(X)
22
23         # 2. Calculate the loss/acc
24         test_loss += loss_fn(test_pred, y)
25         test_acc += accuracy_fn(y_true=y,
26                                y_pred=test_pred.argmax(dim=1)) # go from logits -> prediction labels
27
28 # Adjust metrics and print out
29 test_loss /= len(data_loader)
30 test_acc /= len(data_loader)
31 print(f"Test loss: {test_loss:.5f} | Test acc: {test_acc:.2f}%\n")
32
33 # return the loss value and accuracy for future plots
34 return test_loss.item(), test_acc

1 # Here we are just calling the functions above along with the timer
2 torch.manual_seed(42)
3
4 # Measure time
5 from timeit import default_timer as timer
6 train_time_start_on_gpu = timer()
7
8 # Set epochs
9 epochs = 3
10
11 # Create a optimization and evaluation loop using train_step() and test_step()
12 for epoch in tqdm(range(epochs)):
13     print(f"Epoch: {epoch}\n-----")
14     train_step(model=model_1,
15               data_loader=train_dataloader,
16               loss_fn=loss_fn,
17               optimizer=optimizer,
18               accuracy_fn=accuracy_fn,
19               device=device)
20     test_step(model=model_1,
21              data_loader=test_dataloader,
22              loss_fn=loss_fn,
23              accuracy_fn=accuracy_fn,
24              device=device)

```

```

25
26 train_time_end_on_gpu = timer()
27 total_train_time_model_1 = print_train_time(start=train_time_start_on_gpu,
28                                             end=train_time_end_on_gpu,
29                                             device=device)

```

100% 3/3 [00:18<00:00, 5.97s/it]

Epoch: 0

 Train loss: 3.31834 | Train acc: 7.97%
 Test loss: 3.26682 | Test acc: 11.23%

Epoch: 1

 Train loss: 3.21522 | Train acc: 12.47%
 Test loss: 3.17533 | Test acc: 14.45%

Epoch: 2

 Train loss: 3.11438 | Train acc: 16.30%
 Test loss: 3.09349 | Test acc: 15.14%

Train time on cuda: 18.100 seconds

1 model_0_results

```

{'model_name': 'ASLBaseline',
 'model_loss': 2.8074750900268555,
 'model_acc': 19.43359375}

```

1 # Train time

2 total_train_time_model_0

21.860167875000116

```

1 torch.manual_seed(42)
2 def eval_model(model: torch.nn.Module,
3               data_loader: torch.utils.data.DataLoader,
4               loss_fn: torch.nn.Module,
5               accuracy_fn,
6               device=device):
7     """Returns a dictionary containing the results of model predicting on data_loader."""
8     loss, acc = 0, 0
9     model.eval()
10    with torch.inference_mode():
11        for X, y in tqdm(data_loader):
12            # Make our data device agnostic
13            X, y = X.to(device), y.to(device)
14            # Make predictions
15            y_pred = model(X)
16
17            # Accumulate the loss and acc values per batch
18            loss += loss_fn(y_pred, y)
19            acc += accuracy_fn(y_true=y,
20                             y_pred=y_pred.argmax(dim=1))
21
22    # Scale loss and acc to find the average loss/acc per batch
23    loss /= len(data_loader)
24    acc /= len(data_loader)
25
26    return {"model_name": model.__class__.__name__, # only works when model was created with a class
27          "model_loss": loss.item(),
28          "model_acc": acc}

```

1 # Get model_1 results dictionary

```

2 model_1_results = eval_model(model=model_1,
3                             data_loader=test_data_loader,
4                             loss_fn=loss_fn,
5                             accuracy_fn=accuracy_fn,
6                             device=device)
7 model_1_results

```

```

100% 32/32 [00:01<00:00, 23.82it/s]
{'model_name': 'ASLModelV1',
 'model_loss': 3.0934855937957764,
 'model_acc': 15.13671875}

1 model_0_results

{'model_name': 'ASLBaseline',
 'model_loss': 2.8074750900268555,
 'model_acc': 19.43359375}

```

5 - ASL Model 2: Building a Convolutional Neural Network (CNN)

Now after having implemented a baseline model and a better non-linear model, we've noticed that the accuracy can't even go beyond 50% which is kinda concerning, reason why we are going to implement Convolutional Neural Networks where we would be using two blocks for the data to go through along with 10 hidden units, please find it below

```

1 # Create a convolutional neural network
2 class ASLModelV2(nn.Module):
3     def __init__(self, input_shape: int, hidden_units: int, output_shape: int):
4         super().__init__()
5         self.block_1 = nn.Sequential(
6             nn.Conv2d(in_channels = input_shape,
7                       out_channels = hidden_units,
8                       kernel_size = 3,
9                       stride = 1,
10                      padding = 1),
11             nn.ReLU(),
12             nn.Conv2d(in_channels = hidden_units,
13                      out_channels = hidden_units,
14                      kernel_size = 3,
15                      stride = 1,
16                      padding = 1),
17             nn.ReLU(),
18             nn.MaxPool2d(kernel_size = 2,
19                          stride = 2)
20         )
21         self.block_2 = nn.Sequential(
22             nn.Conv2d(hidden_units, hidden_units, 3, padding = 1),
23             nn.ReLU(),
24             nn.Conv2d(hidden_units, hidden_units, 3, padding = 1),
25             nn.ReLU(),
26             nn.MaxPool2d(2)
27         )
28         self.classifier = nn.Sequential(
29             nn.Flatten(),
30             nn.Linear(in_features = hidden_units * 16 * 16, # used to be 7 * 7
31                      out_features = output_shape)
32         )
33
34     def forward(self, x: torch.Tensor):
35         x = self.block_1(x)
36         #print(f"Output shape of conv_block_1: {x.shape}")
37         x = self.block_2(x)
38         #print(f"Output shape of conv_block_2: {x.shape}")
39         x = self.classifier(x)
40         #print(f"Output shape of classifier: {x.shape}")
41         return x
42
43
44 1 torch.manual_seed(42)
45 2 model_2 = ASLModelV2(input_shape = 3,
46                        hidden_units = 10,
47                        output_shape = len(class_names)).to(device)
48 5 model_2

```

```

ASLModelV2(
  (block_1): Sequential(
    (0): Conv2d(3, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(3): ReLU()
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(block_2): Sequential(
  (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU()
  (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU()
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Flatten(start_dim=1, end_dim=-1)
  (1): Linear(in_features=2560, out_features=29, bias=True)
)
)

```

```

1 # Pass image through model
2 rand_image_tensor = torch.randn(size=(1, 3, 64, 64))
3
4 # Pass image through model
5 model_2(rand_image_tensor.to(device))

```

```

→ tensor([[ 0.1282,  0.0933,  0.0602, -0.0449, -0.0577,  0.0749, -0.0905,  0.1415,
            0.0035,  0.0949, -0.0194, -0.0308,  0.0029, -0.0799,  0.1232,  0.0477,
            0.0469,  0.0493, -0.0039,  0.1222, -0.0345, -0.0842,  0.0605,  0.0132,
            0.0390,  0.1593,  0.0406,  0.1068,  0.0161]], device='cuda:0',
          grad_fn=<AddmmBackward0>)

```

```

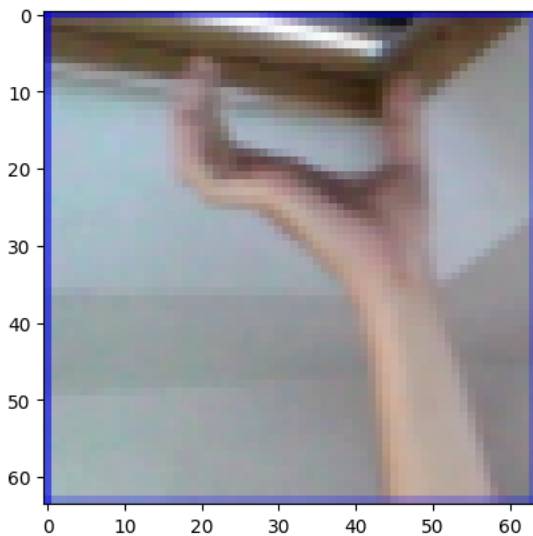
1 image = image * 0.5 + 0.5
2 plt.imshow(image.permute(1, 2, 0))

```

```

→ <matplotlib.image.AxesImage at 0x7e1f9c058690>

```



```

1 model_2.state_dict()

```

```

→ OrderedDict([('block_1.0.weight',
               tensor([[[[ 0.1471,  0.1597, -0.0451],
                           [ 0.1768, -0.0422,  0.0388],
                           [-0.0937,  0.1130,  0.1697]],
                           [[-0.1412,  0.1673,  0.0360],
                           [ 0.1422,  0.0261,  0.0928],
                           [-0.0272,  0.1484,  0.0284]],
                           [[-0.0898,  0.0491, -0.0887],
                           [-0.0226, -0.0782,  0.1277],
                           [-0.1519, -0.0887, -0.0543]]],
                           [[[-0.1157,  0.0182, -0.1901],
                           [ 0.1738, -0.1635,  0.1486],
                           [ 0.0320, -0.0625,  0.1189]],
                           [[ 0.0300,  0.1555,  0.0210],
                           [-0.0607,  0.0517, -0.0522],
                           [ 0.0810,  0.1718,  0.1112]]],

```



```

[[-0.0841,  0.1111,  0.0344],
 [ 0.0977, -0.1173, -0.1905],
 [-0.0744, -0.1476,  0.1579]]],

[[[ 0.0554,  0.0797,  0.0609],
 [-0.0033,  0.1506, -0.1367],
 [ 0.0121, -0.1314,  0.0593]],

 [[-0.0663,  0.0590, -0.0401],
 [ 0.1596, -0.1141, -0.1148],
 [-0.1148,  0.1731,  0.0641]],

 [[ 0.1852, -0.1588, -0.1909],
 [-0.1506, -0.1295,  0.0780],
 [ 0.0689,  0.1599, -0.0994]]],

[[[-0.1312,  0.1021, -0.0778],
 [ 0.1168, -0.0457,  0.1101],
 [-0.1495, -0.0971,  0.0587]],

 [[ 0.0407, -0.0491,  0.1147],
 [ 0.1308, -0.1396, -0.1027],
 [ 0.1762, -0.0649, -0.0682]],

 [[-0.1862, -0.1102,  0.0481],
 [-0.0254, -0.1397,  0.0045],
 [-0.1315, -0.1633, -0.1060]]],

[[[-0.1684, -0.1225,  0.1924],
 [ 0.0363,  0.0593, -0.1795],
 [-0.1264, -0.0641,  0.0301]],

 [-0.1693, -0.0829, -0.1152].

```

5.1 - Stepping through nn.Conv2d()

```

1 torch.manual_seed(42)
2
3 # Create sample batch of random numbers with same size as image batch
4 images = torch.randn(size=(32, 3, 64, 64)) # [batch_size, color_channels, height, width]
5 test_image = images[0] # get a single image for testing
6
7 print(f"Image batch shape: {images.shape} -> [batch_size, color_channels, height, width]")
8 print(f"Single image shape: {test_image.shape} -> [color_channels, height, width]")
9 print(f"Single image pixel values:\n{test_image}")

Image batch shape: torch.Size([32, 3, 64, 64]) -> [batch_size, color_channels, height, width]
Single image shape: torch.Size([3, 64, 64]) -> [color_channels, height, width]
Single image pixel values:
tensor([[[ 1.9269,  1.4873,  0.9007, ...,  1.8446, -1.1845,  1.3835],
          [ 1.4451,  0.8564,  2.2181, ...,  0.3399,  0.7200,  0.4114],
          [ 1.9312,  1.0119, -1.4364, ..., -0.5558,  0.7043,  0.7099],
          ...,
          [-0.5610, -0.4830,  0.4770, ..., -0.2713, -0.9537, -0.6737],
          [ 0.3076, -0.1277,  0.0366, ..., -2.0060,  0.2824, -0.8111],
          [-1.5486,  0.0485, -0.7712, ..., -0.1403,  0.9416, -0.0118]],
        ...,
        [[-0.5197,  1.8524,  1.8365, ...,  0.8935, -1.5114, -0.8515],
          [ 2.0818,  1.0677, -1.4277, ...,  1.6612, -2.6223, -0.4319],
          [-0.1010, -0.4388, -1.9775, ...,  0.2106,  0.2536, -0.7318],
          ...,
          [ 0.2779,  0.7342, -0.3736, ..., -0.4601,  0.1815,  0.1850],
          [ 0.7205, -0.2833,  0.0937, ..., -0.1002, -2.3609,  2.2465],
          [-1.3242, -0.1973,  0.2920, ...,  0.5409,  0.6940,  1.8563]],
        ...,
        [[-0.7978,  1.0261,  1.1465, ...,  1.2134,  0.9354, -0.0780],
          [-1.4647, -1.9571,  0.1017, ..., -1.9986, -0.7409,  0.7011],
          [-1.3938,  0.8466, -1.7191, ..., -1.1867,  0.1320,  0.3407],
          ...,
          [ 0.8206, -0.3745,  1.2499, ..., -0.0676,  0.0385,  0.6335],
          [-0.5589, -0.3393,  0.2347, ...,  2.1181,  2.4569,  1.3083],
          [-0.4092,  1.5199,  0.2401, ..., -0.2558,  0.7870,  0.9924]]])

1 test_image.shape
torch.Size([3, 64, 64])

```

```

1 torch.manual_seed(42)
2
3 # Create a convolutional layer
4 conv_layer = nn.Conv2d(in_channels=3,
5                          out_channels=10,
6                          kernel_size=(3, 3),
7                          stride=1,
8                          padding=0)
9
10 # Pass the data through the convolutional layer
11 conv_output = conv_layer(test_image.unsqueeze(dim=0)) # add an extra dimension for batch
12 conv_output.shape

```

```

→ torch.Size([1, 10, 62, 62])

```

```

1 # Add extra dimension to test image
2 test_image.unsqueeze(dim=0).shape

```

```

→ torch.Size([1, 3, 64, 64])

```

✓ 5.2 - Stepping through nn.MaxPool2d()

```

1 test_image.shape

```

```

→ torch.Size([3, 64, 64])

```

```

1 # Print out original image shape without and with unsqueezed dimension
2 print(f"Test image original shape: {test_image.shape}")
3 print(f"Test image with unsqueezed dimension: {test_image.unsqueeze(dim=0).shape}")
4
5 # Create a sample nn.MaxPool2d() layer
6 max_pool_layer = nn.MaxPool2d(kernel_size=2)
7
8 # Pass data through just the conv_layer
9 test_image_through_conv = conv_layer(test_image.unsqueeze(dim=0))
10 print(f"Shape after going through conv_layer(): {test_image_through_conv.shape}")
11
12 # Pass data through the max pool layer
13 test_image_through_conv_and_max_pool = max_pool_layer(test_image_through_conv)
14 print(f"Shape after going through conv_layer() and max_pool_layer(): {test_image_through_conv_and_max_pool.shape}")

```

```

→ Test image original shape: torch.Size([3, 64, 64])
Test image with unsqueezed dimension: torch.Size([1, 3, 64, 64])
Shape after going through conv_layer(): torch.Size([1, 10, 62, 62])
Shape after going through conv_layer() and max_pool_layer(): torch.Size([1, 10, 31, 31])

```

```

1 torch.manual_seed(42)
2 # Create a random tensor with a similar number of dimensions to our images
3 random_tensor = torch.randn(size=(1, 1, 2, 2))
4 print(f"Random tensor:\n{random_tensor}")
5 print(f"Random tensor shape: {random_tensor.shape}")
6
7 # Create a max pool layer
8 max_pool_layer = nn.MaxPool2d(kernel_size=2) # see what happens when you change the kernel_size value
9
10 # Pass the random tensor through the max pool layer
11 max_pool_tensor = max_pool_layer(random_tensor)
12 print(f"\nMax pool tensor:\n{max_pool_tensor} <- this is the maximum value from random_tensor")
13 print(f"Max pool tensor shape: {max_pool_tensor.shape}")

```

```

→ Random tensor:
tensor([[[[0.3367, 0.1288],
          [0.2345, 0.2303]]]])
Random tensor shape: torch.Size([1, 1, 2, 2])

```

```

Max pool tensor:
tensor([[[[0.3367]]]]) <- this is the maximum value from random_tensor
Max pool tensor shape: torch.Size([1, 1, 1, 1])

```

✓ 5.3 - Setup a loss function and optimizer for model

```

1 # Import accuracy metric
2 from helper_functions import accuracy_fn # Note: could also use torchmetrics.Accuracy(task = 'multiclass', num_classes=len(c
3
4 # Setup loss and optimizer
5 loss_fn = nn.CrossEntropyLoss()
6 optimizer = optim.Adam(model_2.parameters(),
7                           lr = 0.001)
8
9 # Good try :(
10 """loss_fn = nn.CrossEntropyLoss()
11 optimizer = torch.optim.SGD(params=model_2.parameters(),
12                             lr=0.1)"""

```

```

➦ 'loss_fn = nn.CrossEntropyLoss()\noptimizer = torch.optim.SGD(params=model_2.parameters(),\n
0 1)'

```

lr=

✓ 5.4 - Training and Testing model using our training and test functions

```

1 torch.manual_seed(42)
2 torch.cuda.manual_seed(42)
3
4 # Measure time
5 from timeit import default_timer as timer
6 train_time_start_model_2 = timer()
7
8 # Train and test model
9 epochs = 11
10 for epoch in tqdm(range(epochs)):
11     print(f"Epoch: {epoch}\n-----")
12     train_step(data_loader=train_data_loader,
13               model=model_2,
14               loss_fn=loss_fn,
15               optimizer=optimizer,
16               accuracy_fn=accuracy_fn,
17               device=device)
18     test_step(data_loader=test_data_loader,
19              model=model_2,
20              loss_fn=loss_fn,
21              accuracy_fn=accuracy_fn,
22              device=device)
23
24 train_time_end_model_2 = timer()
25 total_train_time_model_2 = print_train_time(start=train_time_start_model_2,
26                                             end=train_time_end_model_2,
27                                             device=device)

```



100%

11/11 [01:10<00:00, 6.25s/it]

Epoch: 0

 Train loss: 3.09838 | Train acc: 12.35%
 Test loss: 2.53919 | Test acc: 26.07%

Epoch: 1

 Train loss: 2.06714 | Train acc: 39.23%
 Test loss: 1.83359 | Test acc: 46.19%

Epoch: 2

 Train loss: 1.35661 | Train acc: 58.58%
 Test loss: 1.55622 | Test acc: 54.69%

Epoch: 3

 Train loss: 0.89766 | Train acc: 71.38%
 Test loss: 1.30300 | Test acc: 63.57%

Epoch: 4

 Train loss: 0.61466 | Train acc: 80.40%
 Test loss: 1.31237 | Test acc: 64.36%

Epoch: 5

 Train loss: 0.39621 | Train acc: 87.28%
 Test loss: 1.40472 | Test acc: 68.36%

Epoch: 6

 Train loss: 0.27861 | Train acc: 91.12%
 Test loss: 1.20559 | Test acc: 73.54%

Epoch: 7

 Train loss: 0.19016 | Train acc: 93.80%
 Test loss: 1.43037 | Test acc: 72.85%

Epoch: 8

 Train loss: 0.14587 | Train acc: 95.40%
 Test loss: 1.44314 | Test acc: 72.27%

Epoch: 9

 Train loss: 0.10936 | Train acc: 96.83%
 Test loss: 1.50674 | Test acc: 73.83%

Epoch: 10

 Train loss: 0.07435 | Train acc: 97.53%
 Test loss: 1.59951 | Test acc: 74.41%

Train time on cuda: 70.527 seconds

```
1 # Get model_2 results
2 model_2_results = eval_model(
3     model=model_2,
4     data_loader=test_dataloader,
5     loss_fn=loss_fn,
6     accuracy_fn=accuracy_fn
7 )
8 model_2_results
```



100%

32/32 [00:01<00:00, 27.43it/s]

```
{'model_name': 'ASLModelV2',
 'model_loss': 1.599514365196228,
 'model_acc': 74.4140625}
```

6 - Compare model results and training time

```

1 import pandas as pd
2 compare_results = pd.DataFrame([model_0_results,
3                                model_1_results,
4                                model_2_results])
5 compare_results

```

	model_name	model_loss	model_acc
0	ASLBaseline	2.807475	19.433594
1	ASLModelV1	3.093486	15.136719
2	ASLModelV2	1.599514	74.414062

```

1 # Add training time to results comparison
2 compare_results["training_time"] = [total_train_time_model_0,
3                                     total_train_time_model_1,
4                                     total_train_time_model_2]
5 compare_results

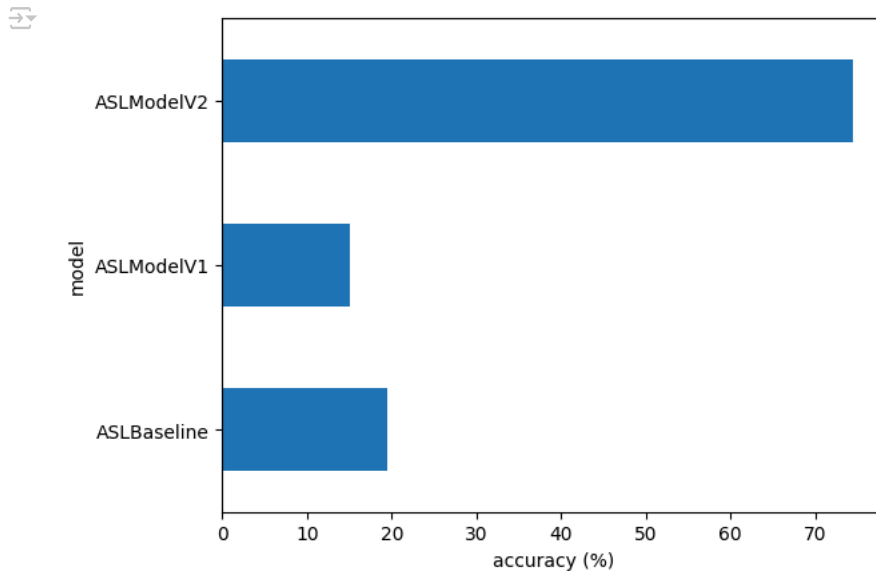
```

	model_name	model_loss	model_acc	training_time
0	ASLBaseline	2.807475	19.433594	21.860168
1	ASLModelV1	3.093486	15.136719	18.099551
2	ASLModelV2	1.599514	74.414062	70.526626

```

1 # Visualize our model results
2 compare_results.set_index("model_name")["model_acc"].plot(kind="barh")
3 plt.xlabel("accuracy (%)")
4 plt.ylabel("model");

```



✓ 7 - Make and evaluate random predictions

```

1 def make_predictions(model: torch.nn.Module,
2                       data: list,
3                       device: torch.device = device):
4     pred_probs = []
5     model.to(device)
6     model.eval()
7     with torch.inference_mode():
8         for sample in data:
9             # Prepare the sample (add a batch dimension and pass to target device)
10            sample = torch.unsqueeze(sample, dim=0).to(device)
11
12            # Forward pass (model outputs raw logits)
13            pred_logit = model(sample)
14
15            # Get prediction probability (logit -> prediction probability)

```

```

16     pred_prob = torch.softmax(pred_logit.squeeze(), dim=0)
17
18     # Get pred_prob off the GPU for further calculations
19     pred_probs.append(pred_prob.cpu())
20
21 # Stack the pred_probs to turn list into a tensor
22 return torch.stack(pred_probs)

```

```

1 import random
2 # random.seed(42)
3 test_samples = []
4 test_labels = []
5 for sample, label in random.sample(list(test_data), k=9):
6     test_samples.append(sample)
7     test_labels.append(label)
8
9 # View the first sample shape
10 test_samples[0].shape

```

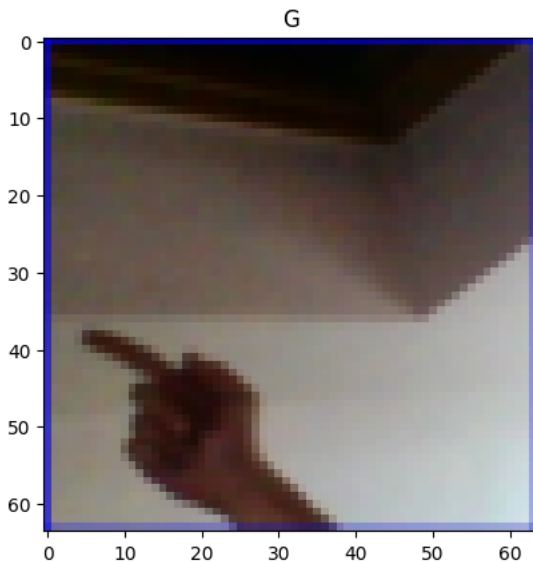
```
→ torch.Size([3, 64, 64])
```

```

1 normalized_test_samples = [sample * 0.5 + 0.5 for sample in test_samples] # Reverse normalization
2 plt.imshow(normalized_test_samples[0].permute(1, 2, 0))
3 plt.title(class_names[test_labels[0]])

```

```
→ Text(0.5, 1.0, 'G')
```



```

1 # Make predictions
2 pred_probs = make_predictions(model=model_2,
3                               data=test_samples)
4
5 # View first two prediction probabilities
6 pred_probs[:2]

```

```
→ tensor([[6.8947e-10, 3.6740e-12, 1.2474e-09, 8.9101e-15, 1.2081e-14, 2.0770e-04,
1.1964e-03, 6.6200e-09, 1.6727e-10, 7.6147e-08, 3.1653e-11, 2.9266e-14,
4.9254e-07, 4.4872e-07, 3.1988e-12, 3.6626e-09, 5.2994e-20, 7.1557e-03,
1.3134e-17, 3.3368e-08, 5.3310e-11, 1.6446e-02, 4.1875e-12, 1.7497e-06,
3.8812e-13, 4.3434e-18, 4.8906e-10, 7.8114e-13, 9.7499e-01],
[9.6114e-13, 4.5881e-17, 3.3778e-12, 1.7567e-10, 2.1434e-08, 2.3279e-10,
3.1673e-10, 1.3308e-10, 4.2215e-07, 2.0996e-05, 6.0409e-09, 5.8643e-07,
2.3985e-02, 2.4080e-03, 2.5756e-06, 2.0818e-06, 8.0176e-06, 8.2397e-05,
1.5705e-05, 4.3257e-05, 1.5224e-05, 4.8218e-03, 8.1809e-02, 7.1317e-04,
8.7580e-04, 4.4739e-05, 5.9217e-02, 3.6400e-13, 8.2593e-01]])
```

```

1 # Convert prediction probabilities to labels
2 pred_classes = pred_probs.argmax(dim=1)
3 pred_classes

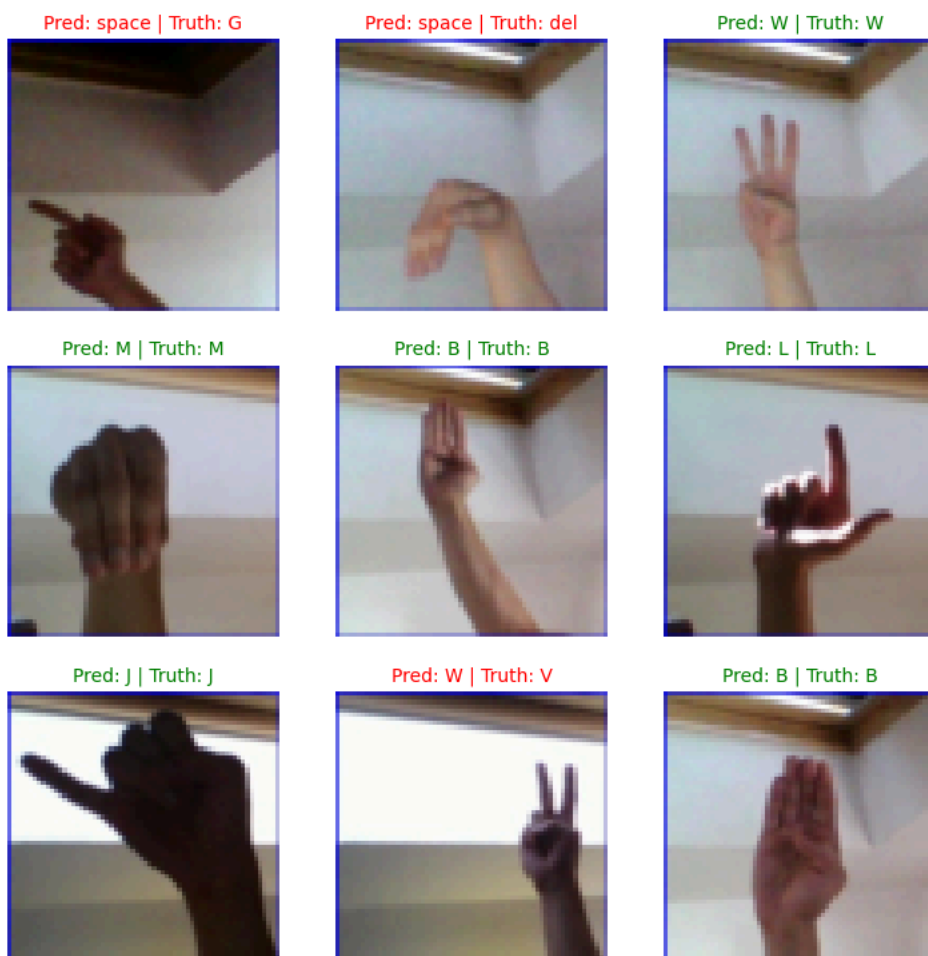
```

```
→ tensor([28, 28, 22, 12,  1, 11,  9, 22,  1])
```

```
1 test_labels
```

```
[6, 26, 22, 12, 1, 11, 9, 21, 1]
```

```
1 # Plot predictions
2 plt.figure(figsize=(9, 9))
3 nrows = 3
4 ncols = 3
5 for i, sample in enumerate(test_samples):
6     # Create subplot
7     plt.subplot(nrows, ncols, i+1)
8
9     normalized_sample = sample * 0.5 + 0.5
10
11     # Plot the target image
12     plt.imshow(normalized_sample.permute(1, 2, 0))
13
14     # Find the prediction (in text form, e.g "Sandal")
15     pred_label = class_names[pred_classes[i]]
16
17     # Get the truth label (in text form)
18     truth_label = class_names[test_labels[i]]
19
20     # Create a title for the plot
21     title_text = f"Pred: {pred_label} | Truth: {truth_label}"
22
23     # Check for equality between pred and truth and change color of title text
24     if pred_label == truth_label:
25         plt.title(title_text, fontsize=10, c="g") # green text if prediction same as truth
26     else:
27         plt.title(title_text, fontsize=10, c="r")
28
29     plt.axis(False);
```



Great! Now we can tell our model is actually learning something and not getting stuck in some accuracy value below 50% :(

✓ 8 - Making a confusion matrix for further prediction evaluation

For those who doesn't know, a confusion matrix is an efficient way of evaluating our classification models by showcasing data visualizations where we can tell how accurate every class is based on the number of sample each one of them have

```
1 # Import tqdm.auto
2 from tqdm.auto import tqdm
3
4
5 # 1. Make predictions with trained model
6 y_preds = []
7 model_2.eval()
8 with torch.inference_mode():
9     for X, y in tqdm(test_dataloader, desc="Making predictions.."):
10         # Send the data and targets to target device
11         X, y = X.to(device), y.to(device)
12         # Do the forward pass
13         y_logits = model_2(X)
14         # Turn predictions from logits -> prediction probabilities -> prediction labels
15         y_pred = torch.softmax(y_logits.squeeze(), dim=0).argmax(dim=1)
16         # Put prediction on CPU for evaluation
17         y_preds.append(y_pred.cpu())
18
19 # Concatenate list of predictions into a tensor
20 y_pred_tensor = torch.cat(y_preds)
21 y_pred_tensor
```




Making predictions...: 100%

32/32 [00:01<00:00, 25.45it/s]

```

tensor([ 5, 19,  3, 24, 10, 13,  1, 17, 26,  3, 12, 22,  4, 24, 14,  0, 11,  3,
        27, 19, 13,  2,  0,  9,  4,  7, 20, 14, 14, 17, 12,  8,  2,  7, 14,  3,
        17, 17, 23, 25, 15, 20, 13,  8,  5,  4,  6, 19, 20,  7, 22, 23, 24, 14,
        19, 12, 28, 19, 21, 16, 13, 28, 26, 21, 26,  6, 28, 25,  2,  4, 22, 21,
        8, 10, 19, 20, 13, 11,  2,  9, 23, 19,  1, 10, 16, 24, 27, 12, 12, 21,
        6, 10, 24,  8,  5, 20, 23, 27, 22, 26,  9, 12,  0,  9,  3,  2,  4,  6,
        24, 21,  9, 28, 13, 19, 22, 22,  8, 22, 16, 15,  1, 24, 17, 11, 24, 24,
        7, 11, 10,  6,  0, 14,  5,  4, 10,  2, 25, 26,  9, 23, 17, 16, 27, 26,
        3,  1,  8, 28, 13,  6, 19,  7, 12,  6,  6, 15, 20, 22,  5,  9, 20, 12,
        4,  8, 21,  0, 24, 21, 24, 12, 28, 15,  9, 15,  2, 12,  4,  9, 12, 27,
        22,  8,  2,  2, 23, 21,  7, 14, 13, 28, 24, 12, 26,  8, 23, 24,  6, 27,
        12, 20, 24, 10, 23,  4, 21,  7, 15,  6,  5, 16, 14, 22,  2, 10,  9,  5,
        25, 19, 11, 28, 17, 19,  1, 18,  0, 16,  3,  8,  7, 22, 15, 14,  2,  8,
        7,  8, 10, 25, 19, 12,  8, 16,  6, 13, 22, 17, 28,  5, 21, 27,  1, 22,
        9, 15, 12, 26,  7,  9,  1, 10, 25, 12, 23,  8,  0, 28, 13, 14, 22,  4,
        18, 16,  5, 15, 24, 14, 12, 28, 23, 15, 15, 27, 20, 11, 21,  3,  2, 28,
        20, 28, 28,  0, 19, 20,  9, 13,  8,  6,  3, 10,  2, 21,  2, 14, 14, 25,
        18, 16, 23, 11, 18, 15, 11,  3, 26, 27,  1, 11, 28, 14,  9, 13,  3,  8,
        11, 16, 21,  7, 19, 11, 22, 28,  4, 19, 10, 25, 27, 26, 28,  6, 23, 23,
        9, 15, 25, 24, 10, 17, 28, 26, 26, 23,  1, 26, 23, 11, 11, 19,  2, 17,
        23, 21, 23, 12, 19, 26, 10,  9, 17, 12, 16, 20, 18, 28, 22, 25, 12, 25,
        12,  9,  7,  4,  5, 15, 20, 21, 27, 10, 14,  1, 12, 15, 10, 28,  8,  3,
        13, 18,  6, 23, 14, 19, 24,  5, 25, 11,  6,  2,  4,  5, 20, 20,  5,  4,
        12,  2,  3, 27, 25, 21,  0, 21,  6,  5, 15, 12, 19,  9,  2, 24, 20, 26,
        9, 19, 12, 12,  9, 23, 18,  7, 10, 28, 23,  2, 17, 24, 24,  8, 11,  0,
        21, 15, 16,  1, 28, 22, 10, 16, 12, 24, 25,  8,  0, 17,  6, 28, 19,  9,
        28, 23, 12, 26, 21,  3, 18, 23, 10, 20,  5,  2, 16,  2,  8, 15, 20, 22,
        18,  2,  0, 11,  1,  7, 24,  2, 28, 12, 26, 11, 16,  5, 14,  6,  9, 21,
        4,  8,  1, 18, 11, 27, 17, 19, 13, 28, 24, 18, 15, 25, 12, 23, 13, 19,
        0,  7, 17,  8, 28, 28, 11,  2, 20, 26, 18,  3, 23, 12,  8, 23,  5, 14,
        1, 10, 21, 10, 22, 21,  2, 23, 13,  1,  0, 26,  9,  5, 17, 14, 13, 14,
        19,  2, 25, 12, 15, 19,  3, 23, 20, 27, 16, 24, 14,  5,  6, 25, 11, 26,
        7,  4,  5, 21,  4, 23, 16, 12, 19, 13, 22,  6, 11,  2, 25,  5, 12, 27,
        10, 12,  6, 11, 26, 17, 15, 11, 27,  6, 24,  5, 28,  8, 26,  3, 13, 19,
        14, 11, 26, 22,  6, 26, 24,  4, 28,  5,  5, 20, 16,  2, 12,  1, 19, 22,
        19, 21, 11, 20,  8, 17, 15, 18, 11, 22, 19,  9,  1, 12,  6, 22, 19, 24,
        14, 20, 21, 25, 12, 23,  7, 10, 13, 19,  6, 26,  3,  8, 17, 10, 19, 14,
        20, 15,  6,  5,  9, 12,  9,  1,  0, 19, 25,  4, 15,  6, 26, 17, 20,  2,
        26, 17, 14, 11, 12, 23,  5,  8,  5, 27, 10, 17, 15, 19, 21,  0, 13,  8,
        14, 16, 14, 23, 28,  6, 23,  7, 22,  9,  2, 13, 14,  1,  8, 18, 12,  4,
        17, 19, 16, 10, 15, 15, 19, 23, 20, 22, 25, 24, 14, 21, 23,  5, 14, 13,
        5, 20,  9, 15, 19,  6,  5, 11, 18, 11, 13,  1,  8, 19, 28, 10,  0,  5,
        14, 21, 17, 12,  3, 28,  4, 14,  5, 13,  5, 24, 11,  0, 13, 22,  7,  9,
        21, 28, 15, 23, 25, 20,  3,  2,  8, 28, 25, 23, 17, 10, 18, 12,  5, 28,
        18, 27,  9, 28, 26, 17,  1, 24, 25, 21, 25, 27,  2, 12, 22, 26, 23, 28,
        4, 14, 24,  1, 24, 12,  7, 20,  1, 20, 17, 24, 24, 28, 13,  7, 11, 14,
        20, 25,  3, 19, 10, 23, 20, 28,  2, 24, 21, 19, 23, 16,  8,  7, 23,  9,
        28,  1,  6, 12, 11, 24, 28, 28,  0, 18, 15, 22,  8, 12, 13, 17, 18, 26,
        9, 13,  1,  6, 17, 10,  5, 22,  9, 23, 16,  6, 26, 15, 23, 11, 27, 19,
        20, 28, 13,  0, 12, 15, 21,  0, 24, 26, 19,  5,  3,  8,  8,  9, 14, 12,
        26, 26, 14, 11,  1,  5,  0, 13, 18,  7, 23, 28,  2,  8, 15, 22, 10, 15,
        23, 21, 17, 14, 25, 12, 19, 16, 22, 20, 28, 12, 13, 13, 20,  8, 17,  4,
        21, 15, 13,  6, 10, 10,  1,  6,  5,  5, 28, 17, 14, 24, 25, 12, 25,  2,
        27,  9, 26, 19, 11, 25,  0, 25, 17, 23, 20,  6, 18, 12, 24,  1,  8, 13,
        21, 26, 25, 17, 25,  4,  9, 28,  2, 13, 19, 12,  1, 10, 25,  1, 23, 20,
        27, 14, 19,  3, 18,  8, 24, 15, 22,  1])

```

```
1 len(y_pred_tensor)
```



1000

```

1 # See if required packages are installed and if not, install them...
2 try:
3     import torchmetrics, mlxtend
4     print(f"mlxtend version: {mlxtend.__version__}")
5     assert int(mlxtend.__version__.split(".")[1]) >= 19, "mlxtend version should be 0.19.0 or higher"
6 except:
7     !pip install torchmetrics -U mlxtend
8     import torchmetrics, mlxtend
9     print(f"mlxtend version: {mlxtend.__version__}")

```



```

mlxtend version: 0.23.4
Requirement already satisfied: torchmetrics in /usr/local/lib/python3.11/dist-packages (1.7.1)
Requirement already satisfied: mlxtend in /usr/local/lib/python3.11/dist-packages (0.23.4)
Requirement already satisfied: numpy>1.20.0 in /usr/local/lib/python3.11/dist-packages (from torchmetrics) (2.0.2)
Requirement already satisfied: packaging>17.1 in /usr/local/lib/python3.11/dist-packages (from torchmetrics) (24.2)
Requirement already satisfied: torch>=2.0.0 in /usr/local/lib/python3.11/dist-packages (from torchmetrics) (2.6.0+cu124)
Requirement already satisfied: lightning-utilities>=0.8.0 in /usr/local/lib/python3.11/dist-packages (from torchmetrics) (0.

```

```

Requirement already satisfied: scipy>=1.2.1 in /usr/local/lib/python3.11/dist-packages (from mlxtend) (1.14.1)
Requirement already satisfied: pandas>=0.24.2 in /usr/local/lib/python3.11/dist-packages (from mlxtend) (2.2.2)
Requirement already satisfied: scikit-learn>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from mlxtend) (1.6.1)
Requirement already satisfied: matplotlib>=3.0.0 in /usr/local/lib/python3.11/dist-packages (from mlxtend) (3.10.0)
Requirement already satisfied: joblib>=0.13.2 in /usr/local/lib/python3.11/dist-packages (from mlxtend) (1.4.2)
Requirement already satisfied: setuptools in /usr/local/lib/python3.11/dist-packages (from lightning-utilities>=0.8.0->torch)
Requirement already satisfied: typing_extensions in /usr/local/lib/python3.11/dist-packages (from lightning-utilities>=0.8.0->torch)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend) (11.1.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.11/dist-packages (from matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24.2->mlxtend) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24.2->mlxtend) (2025.1)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.11/dist-packages (from scikit-learn>=1.3.1->mlxtend)
Requirement already satisfied: filelock in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (3.18.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (3.4.2)
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (3.1.6)
Requirement already satisfied: fsspec in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (2025.3.2)
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cudnn-cu12==9.1.0.70 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cublas-cu12==12.4.5.8 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cufft-cu12==11.2.1.3 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-curand-cu12==10.3.5.147 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cusolver-cu12==11.6.1.9 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cusparse-cu12==12.3.1.170 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.2 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-nccl-cu12==2.21.5 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-nvtx-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.4.127 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics)
Requirement already satisfied: triton==3.2.0 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (3.2.0)
Requirement already satisfied: sympy==1.13.1 in /usr/local/lib/python3.11/dist-packages (from torch>=2.0.0->torchmetrics) (1.13.1)
Requirement already satisfied: mpmath<1.4, >=1.1.0 in /usr/local/lib/python3.11/dist-packages (from sympy==1.13.1->torch>=2.0.0->torchmetrics)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.7->matplotlib>=3.0.0->mlxtend)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.11/dist-packages (from Jinja2->torch>=2.0.0->torchmetrics)
mlxtend version: 0.23.4

```

```

1 import mlxtend
2 print(mlxtend.__version__)

```

0.23.4

```
1 class_names
```

```

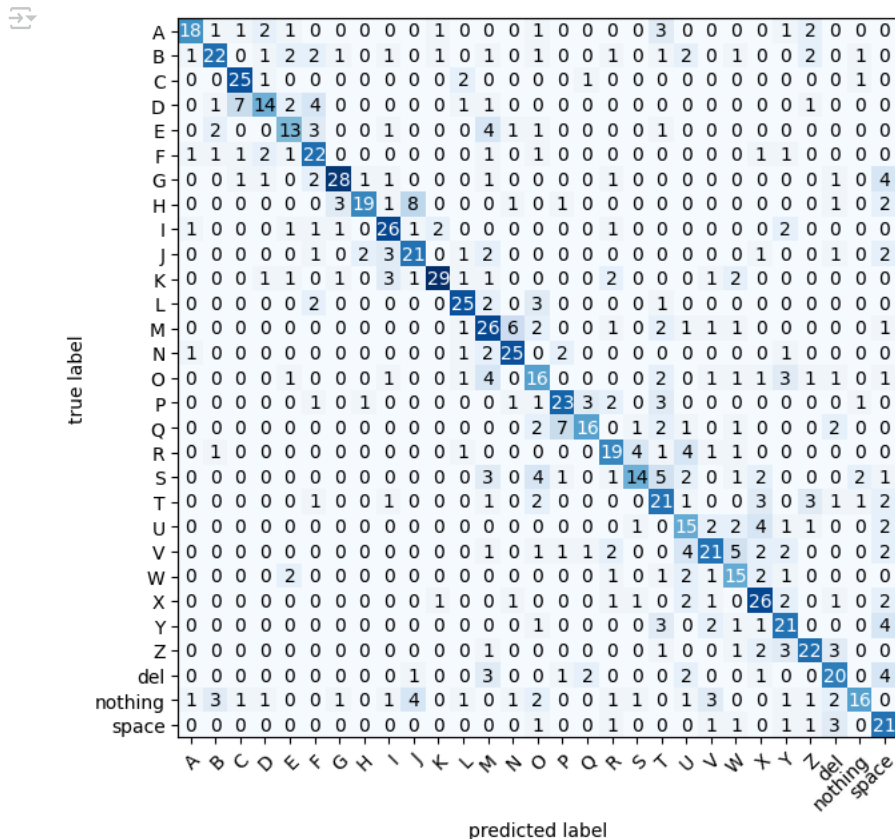
['A',
 'B',
 'C',
 'D',
 'E',
 'F',
 'G',
 'H',
 'I',
 'J',
 'K',
 'L',
 'M',
 'N',
 'O',
 'P',
 'Q',
 'R',
 'S',
 'T',
 'U',
 'V',
 'W',
 'X',
 'Y',
 'Z',
 'del',
 'nothing',
 'space']

```

```
1 y_pred_tensor[:10]
```

```
tensor([ 5, 19,  3, 24, 10, 13,  1, 17, 26,  3])
```

```
1 from torchmetrics import ConfusionMatrix
2 from mlxtend.plotting import plot_confusion_matrix
3
4 # 2. Setup confusion instance and compare predictions to targets
5 confmat = ConfusionMatrix(task="multiclass", num_classes=len(class_names))
6
7 # Here we are using a list comprehension that extracts the target labels from our test data,
8 # in other words, this is just iterating through our test dataset taking out only the target labels
9 # and putting them into a new list called targets in this case
10 targets = [label for _, label in test_data]
11
12 confmat_tensor = confmat(preds=y_pred_tensor,
13                          target=torch.tensor(targets))
14                          #target=[label for _, label in test_data])
15                          #target=test_data.targets)
16
17 # 3. Plot the confusion matrix
18 fig, ax = plot_confusion_matrix(
19     conf_mat=confmat_tensor.numpy(), # matplotlib likes working with numpy
20     class_names=class_names,
21     figsize=(10, 7)
22 )
```



9 - Save and Load Model

```
1 from pathlib import Path
2
3 # Create model dictory path
4 MODEL_PATH = Path("models")
5 MODEL_PATH.mkdir(parents=True,
6                  exist_ok=True)
7
8 # Create model save
9 MODEL_NAME = "ASL_Model.pth"
10 MODEL_SAVE_PATH = MODEL_PATH / MODEL_NAME
11
12 # Save the model state dict
```

```

13 print(f"Saving model to: {MODEL_SAVE_PATH}")
14 torch.save(obj=model_2.state_dict(),
15             f=MODEL_SAVE_PATH)

```

→ Saving model to: models/ASL_Model.pth

```

1 image_shape = [3, 64, 64]

1 # Create a new instance
2 torch.manual_seed(42)
3
4 loaded_model_2 = ASLModelV2(input_shape=3,
5                             hidden_units=10,
6                             output_shape=len(class_names))
7
8 # Load in the save state_dict()
9 loaded_model_2.load_state_dict(torch.load(f=MODEL_SAVE_PATH))
10
11 # Send the model to the target device
12 loaded_model_2.to(device)

```

→ ASLModelV2(
 (block_1): Sequential(
 (0): Conv2d(3, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (1): ReLU()
 (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (3): ReLU()
 (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
 (block_2): Sequential(
 (0): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (1): ReLU()
 (2): Conv2d(10, 10, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
 (3): ReLU()
 (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
 (classifier): Sequential(
 (0): Flatten(start_dim=1, end_dim=-1)
 (1): Linear(in_features=2560, out_features=29, bias=True)
)
)

```

1 # Evaluate loaded model
2 torch.manual_seed(42)
3
4 loaded_model_2_results = eval_model(
5     model=loaded_model_2,
6     data_loader=test_data_loader,
7     loss_fn=loss_fn,
8     accuracy_fn=accuracy_fn
9 )
10
11 loaded_model_2_results

```

→ 100% 32/32 [00:01<00:00, 27.12it/s]

```

{'model_name': 'ASLModelV2',
 'model_loss': 1.599514365196228,
 'model_acc': 74.4140625}

```

```
1 model_2_results
```

→ {'model_name': 'ASLModelV2',
 'model_loss': 1.599514365196228,
 'model_acc': 74.4140625}

```

1 # Check if model results are close to each other
2 torch.isclose(torch.tensor(model_2_results["model_loss"]),
3               torch.tensor(loaded_model_2_results["model_loss"]),
4               atol=1e-02)

```

→ tensor(True)

✓ 10 - Transfer Learning

After having started a baseline model, then building a better model with non-linearity, and then a CNN, we decided to perform transfer learning using a pre-trained ResNet18 model to find the most feasible model among the ones we have so far. We need a model that has the highest efficiency and accuracy and the lowest running time and loss.

We resized images to 200x200(same as originals) to hold the quality during training, converted them to tensors, and normalized pixel values to a common scale.

We loaded the dataset using ImageFolder, and increment the subset by using 10,000 samples (we need to find a balance between image size and number of samples to get the best performance). We decided to increment samples and size because ResNet18 allows us to perform faster computations and is used to process high amount of data.

We split the subset into 80% training and 20% testing like before.

Then we loaded the ResNet18 model and we decided to froze its convolutional layers so that it can retain the visual features it has already learned from ImageNet.

Just the final fully-connected (fc) layer is replaced and trained. We believe this reduced the training time and the risk of overfitting (that we have before) specially with this new dataset.

We used helper function like train_step and test_step to to track their results after each epoch.

After the training, we visualize the training and test loss curves, to see how well the model is learning. Then we printed the predicted and actual labels to see how the model performs after the training and to check if everything is okay.

To conclude we find transfer learning to be the most efficient among the models that we built, demonstrating less run time and loss, and high accuracy.

Resnet18 documentation: <https://debuggercafe.com/implementing-resnet18-in-pytorch-from-scratch/>

✓ 10.1 - Create DataLoader for New Dataset

```

1 from torchvision import datasets, transforms
2 from torch.utils.data import DataLoader
3 #from torchinfo import summary
4
5 device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
6
7 torch.manual_seed(42)
8 random.seed(42)
9
10 # Path to your new dataset
11 new_data_dir = "/content/asl_alphabet_data/asl_alphabet_train/asl_alphabet_train"
12
13 """transforms.Normalize([0.485, 0.456, 0.406],
14                          [0.229, 0.224, 0.225])"""
15
16 # Define transforms
17 new_transform = transforms.Compose([
18     transforms.Resize((200, 200)),
19     transforms.ToTensor(),
20     transforms.Normalize([0.5, 0.5, 0.5], [0.5, 0.5, 0.5])
21 ])
22
23 # Create the dataset
24 new_dataset = datasets.ImageFolder(root=new_data_dir, transform=new_transform)
25
26 # Use a smaller subset (10000 samples)
27 subset_indices = random.sample(range(len(new_dataset)), 10000)
28 subset = Subset(new_dataset, subset_indices)
29
30 train_size = int(0.8 * len(subset))
31 test_size = len(subset) - train_size
32 train_dataset, test_dataset = random_split(subset, [train_size, test_size])
33
34
35 # Create the DataLoader
36 train_dataloader = DataLoader(train_dataset,
37                               batch_size=32,
38                               shuffle=True)
39 test_dataloader = DataLoader(test_dataset,
40                              batch_size=32,
41                              shuffle=False)

```

```

42
43 num_new_classes = len(new_dataset.classes)
44

```

✓ 10.2 - Freeze Layers and changing the output layer to suit our needs

```

1 base_model = models.resnet18(pretrained=True)
2
3 for param in base_model.parameters():
4     param.requires_grad = False
5
6 base_model.fc = nn.Linear(base_model.fc.in_features, num_new_classes)
7 base_model = base_model.to(device)

```

 /usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated. Please use 'weights' instead.
 warnings.warn(
 /usr/local/lib/python3.11/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or
 warnings.warn(msg)

```

1 len(class_names)

```

 29

✓ 10.3 - Train Model

```


1 # Define loss function and optimizer
2 loss_fn = nn.CrossEntropyLoss()
3 optimizer = torch.optim.Adam(base_model.fc.parameters(), lr=0.001)

```

```

1 # Lists to store losses for plotting
2 train_losses = []
3 test_losses = []
4 train_accuracies = []
5 test_accuracies = []
6
7 # Training loop
8 epochs = 10
9 for epoch in range(epochs):
10     print(f"Epoch: {epoch}\n-----")
11
12     # Train on the new dataset and get train loss
13     train_loss, train_acc = train_step(data_loader=train_dataloader,
14                                       model= base_model,
15                                       loss_fn=loss_fn,
16                                       optimizer=optimizer,
17                                       accuracy_fn=accuracy_fn,
18                                       device=device)
19     train_losses.append(train_loss) # Append train loss and accuracy to list
20     train_accuracies.append(train_acc)
21
22     # Evaluate on the test dataset and get test loss
23     test_loss, test_acc = test_step(data_loader=test_dataloader,
24                                   model= base_model,
25                                   loss_fn=loss_fn,
26                                   accuracy_fn=accuracy_fn,
27                                   device=device)
28     test_losses.append(test_loss) # Append test loss and accuracy to list
29     test_accuracies.append(test_acc)

```

 Epoch: 0

Train loss: 2.20810 | Train acc: 47.80%

Test loss: 1.40413 | Test acc: 73.71%

Epoch: 1

Train loss: 1.12713 | Train acc: 78.04%

Test loss: 0.92178 | Test acc: 80.70%

Epoch: 2

Train loss: 0.79570 | Train acc: 84.03%

Test loss: 0.71868 | Test acc: 84.47%

Epoch: 3

Train loss: 0.63709 | Train acc: 86.97%
Test loss: 0.59842 | Test acc: 85.81%

Epoch: 4

Train loss: 0.53795 | Train acc: 88.91%
Test loss: 0.52690 | Test acc: 88.05%

Epoch: 5

Train loss: 0.46431 | Train acc: 90.24%
Test loss: 0.48939 | Test acc: 87.70%

Epoch: 6

Train loss: 0.40846 | Train acc: 91.54%
Test loss: 0.45753 | Test acc: 88.69%

Epoch: 7

Train loss: 0.36935 | Train acc: 92.39%
Test loss: 0.40096 | Test acc: 89.48%

Epoch: 8

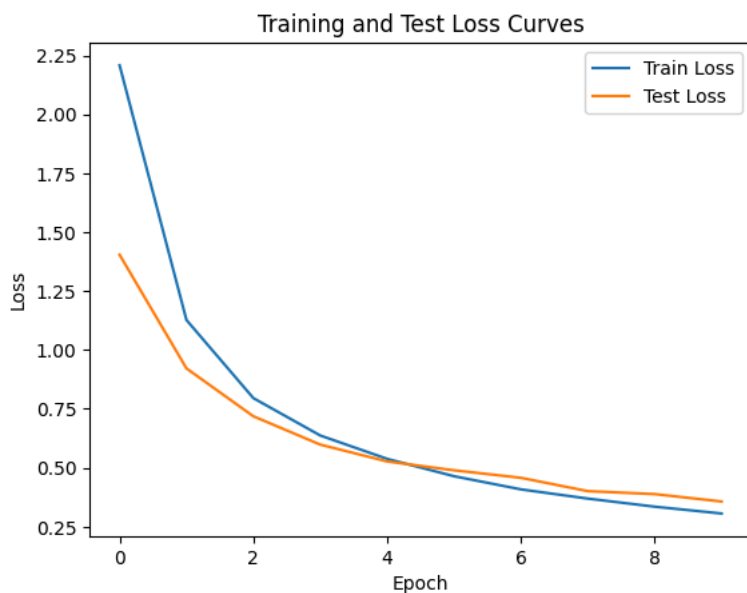
Train loss: 0.33506 | Train acc: 92.71%
Test loss: 0.38843 | Test acc: 89.58%

Epoch: 9

Train loss: 0.30606 | Train acc: 93.50%
Test loss: 0.35697 | Test acc: 90.72%

✓ 10.4 - Evaluate model by plotting loss curves

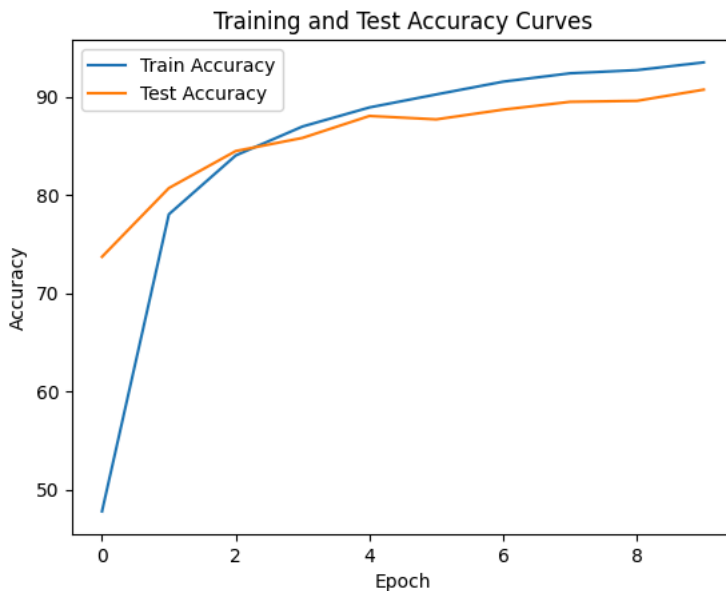
```
1 import matplotlib.pyplot as plt
2
3 # Plot the losses
4 plt.plot(train_losses, label='Train Loss')
5 plt.plot(test_losses, label='Test Loss')
6 plt.title('Training and Test Loss Curves')
7 plt.xlabel('Epoch')
8 plt.ylabel('Loss')
9 plt.legend()
10 plt.show()
```



```

1 # Plot the accuracies
2 plt.plot(train_accuracies, label='Train Accuracy')
3 plt.plot(test_accuracies, label='Test Accuracy')
4 plt.title('Training and Test Accuracy Curves')
5 plt.xlabel('Epoch')
6 plt.ylabel('Accuracy')
7 plt.legend()
8 plt.show()

```



✓ 10.5 - Transfer learning confusion matrix

```

1 # Import tqdm.auto
2 from tqdm.auto import tqdm
3
4 # 1. Make predictions with trained model
5 y_preds = []
6 model_2.eval()
7 with torch.inference_mode():
8     for X, y in tqdm(test_dataloader, desc="Making predictions..."):
9         # Send the data and targets to target device
10        X, y = X.to(device), y.to(device)
11        # Do the forward pass
12        y_logits = base_model(X)
13        # Turn predictions from logits -> prediction probabilities -> prediction labels
14        y_pred = torch.softmax(y_logits.squeeze(), dim=0).argmax(dim=1)
15        # Put prediction on CPU for evaluation
16        y_preds.append(y_pred.cpu())
17
18 # Concatenate list of predictions into a tensor
19 y_pred_tensor = torch.cat(y_preds)
20 y_pred_tensor
21
22
23 from torchmetrics import ConfusionMatrix
24 from mlxtend.plotting import plot_confusion_matrix
25
26 # 2. Setup confusion instance and compare predictions to targets
27 confmat = ConfusionMatrix(task="multiclass", num_classes=len(class_names))
28
29 # Here we are using a list comprehension that extracts the target labels from our test data,
30 # in other words, this is just iterating through our test dataset taking out only the target label
31 # and putting them into a new list called targets in this case
32 targets = []
33 for _, label in test_dataloader: # iterate through the test_dataloader to get all target labels
34     targets.extend(label.tolist()) # add the labels from the current batch to the targets list
35
36 confmat_tensor = confmat(preds=y_pred_tensor,
37                           target=torch.tensor(targets))
38 #target=[label for _, label in test_data]
39 #target=test_data.targets)

```